

# PROJECT NAME: AI BASED DIABETES PREDICTION SYSTEM.

**PHASE-04 : developing project part 2**

**Mail id:nandhiniselvam0405@gmail.com**

In terms of data integration, the system leverages not only traditional patient data sources but also incorporates wearable device data, continuous glucose monitoring, and real-time health metrics. This expanded dataset enables a more comprehensive understanding of an individual's health status and lifestyle, resulting in more precise predictions. The predictive modeling component incorporates state-of-the-art machine learning algorithms, deep learning techniques, and artificial intelligence to create highly accurate prediction models. These models adapt and improve over time by continuously learning from new data, allowing for dynamic updates and personalized predictions.

**Machine Learning and AI:** Machine learning algorithms, especially deep learning, have shown promise in predicting diabetes. These algorithms can analyze large datasets of patient information, such as medical records, genetic data, and lifestyle factors, to identify patterns and predict the risk of diabetes.

**Artificial Pancreas Systems:** These systems combine insulin pumps and CGM devices with predictive algorithms to automate insulin delivery. They can predict future glucose levels and adjust insulin delivery accordingly, reducing the risk of hypoglycemia and hyperglycemia.


**Mobile Apps and Wearables:** There's a growing ecosystem of mobile apps and wearables designed to help individuals manage diabetes. These apps often include predictive features that use data on diet, exercise, and glucose levels to provide personalized recommendations and forecasts.

```
import numpy as np
import pandas as pd
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn import svm
from sklearn.metrics import accuracy_score
```

```
# loading the dataset to the pandas dataframe
diabetes_dataset = pd.read_csv('/content/diabetes.csv')
```

```
pd.read_csv?
```

```
# printing the first 5 rows of the dataset
diabetes_dataset.head()
```



|   | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI  | DiabetesPedigreeFunction | Age | Outcome |
|---|-------------|---------|---------------|---------------|---------|------|--------------------------|-----|---------|
| 0 | 6           | 148     | 72            | 35            | 0       | 33.6 | 0.627                    | 50  | 1       |
| 1 | 1           | 85      | 66            | 29            | 0       | 26.6 | 0.351                    | 31  | 0       |
| 2 | 8           | 183     | 64            | 0             | 0       | 23.3 | 0.672                    | 32  | 1       |
| 3 | 1           | 89      | 66            | 23            | 94      | 28.1 | 0.167                    | 21  | 0       |
| 4 | 0           | 137     | 40            | 35            | 168     | 43.1 | 2.288                    | 33  | 1       |

```
# number of rows and column in this dataset
diabetes_dataset.shape
```

```
(768, 9)
```

```
# getting the statistical measures of the data
diabetes_dataset.describe()
```

|       | Pregnancies | Glucose    | BloodPressure | SkinThickness | Insulin    | BMI        |
|-------|-------------|------------|---------------|---------------|------------|------------|
| count | 768.000000  | 768.000000 | 768.000000    | 768.000000    | 768.000000 | 768.000000 |
| mean  | 3.845052    | 120.894531 | 69.105469     | 20.536458     | 79.799479  | 31.992578  |
| std   | 3.369578    | 31.972618  | 19.355807     | 15.952218     | 115.244002 | 7.884160   |
| min   | 0.000000    | 0.000000   | 0.000000      | 0.000000      | 0.000000   | 0.000000   |
| 25%   | 1.000000    | 99.000000  | 62.000000     | 0.000000      | 0.000000   | 27.300000  |
| 50%   | 3.000000    | 117.000000 | 72.000000     | 23.000000     | 30.500000  | 32.000000  |
| 75%   | 6.000000    | 140.250000 | 80.000000     | 32.000000     | 127.250000 | 36.600000  |
| max   | 17.000000   | 199.000000 | 122.000000    | 99.000000     | 846.000000 | 67.100000  |

```
diabetes_dataset['Outcome'].value_counts()
```

```
0    500
1    268
Name: Outcome, dtype: int64
```

0--> Non-Diabetic

1--> Diabetic

```
diabetes_dataset.groupby('Outcome').mean()
```

|         | Pregnancies | Glucose    | BloodPressure | SkinThickness | Insulin    | BMI       |
|---------|-------------|------------|---------------|---------------|------------|-----------|
| Outcome |             |            |               |               |            |           |
| 0       | 3.298000    | 109.980000 | 68.184000     | 19.664000     | 68.792000  | 30.304200 |
| 1       | 4.865672    | 141.257463 | 70.824627     | 22.164179     | 100.335821 | 35.142537 |

```
# seperating the data and labels
X = diabetes_dataset.drop(columns = 'Outcome', axis=1)
Y = diabetes_dataset['Outcome']
```

```
print(X)
```

|     | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI  | \ |
|-----|-------------|---------|---------------|---------------|---------|------|---|
| 0   | 6           | 148     | 72            | 35            | 0       | 33.6 |   |
| 1   | 1           | 85      | 66            | 29            | 0       | 26.6 |   |
| 2   | 8           | 183     | 64            | 0             | 0       | 23.3 |   |
| 3   | 1           | 89      | 66            | 23            | 94      | 28.1 |   |
| 4   | 0           | 137     | 40            | 35            | 168     | 43.1 |   |
| ..  | ...         | ...     | ...           | ...           | ...     | ...  |   |
| 763 | 10          | 101     | 76            | 48            | 180     | 32.9 |   |
| 764 | 2           | 122     | 70            | 27            | 0       | 36.8 |   |
| 765 | 5           | 121     | 72            | 23            | 112     | 26.2 |   |
| 766 | 1           | 126     | 60            | 0             | 0       | 30.1 |   |
| 767 | 1           | 93      | 70            | 31            | 0       | 30.4 |   |

|     | DiabetesPedigreeFunction | Age |
|-----|--------------------------|-----|
| 0   | 0.627                    | 50  |
| 1   | 0.351                    | 31  |
| 2   | 0.672                    | 32  |
| 3   | 0.167                    | 21  |
| 4   | 2.288                    | 33  |
| ..  | ...                      | ... |
| 763 | 0.171                    | 63  |
| 764 | 0.340                    | 27  |
| 765 | 0.245                    | 30  |
| 766 | 0.349                    | 47  |
| 767 | 0.315                    | 23  |

[768 rows x 8 columns]

```
print(Y)
```

```
0      1
1      0
2      1
3      0
4      1
..
763    0
764    0
765    0
766    1
767    0
Name: Outcome, Length: 768, dtype: int64
```

## Data Standardization

```
scaler = StandardScaler()
```

```
scaler.fit(X)
```

```
▼ StandardScaler
StandardScaler()
```

```
standardized_data = scaler.transform(X)
```

```
print(standardized_data)
```

```
[[ 0.63994726  0.84832379  0.14964075 ...  0.20401277  0.46849198
   1.4259954 ]
 [-0.84488505 -1.12339636 -0.16054575 ... -0.68442195 -0.36506078
  -0.19067191]
 [ 1.23388019  1.94372388 -0.26394125 ... -1.10325546  0.60439732
  -0.10558415]
 ...
 [ 0.3429808  0.00330087  0.14964075 ... -0.73518964 -0.68519336
  -0.27575966]
 [-0.84488505  0.1597866  -0.47073225 ... -0.24020459 -0.37110101
   1.17073215]
 [-0.84488505 -0.8730192  0.04624525 ... -0.20212881 -0.47378505
  -0.87137393]]
```

```
X = standardized_data
Y = diabetes_dataset['Outcome']
```

```
print(X)
print(Y)
```

```
[[ 0.63994726  0.84832379  0.14964075 ...  0.20401277  0.46849198
   1.4259954 ]
 [-0.84488505 -1.12339636 -0.16054575 ... -0.68442195 -0.36506078
  -0.19067191]]
```

```
[ 1.23388019  1.94372388 -0.26394125 ... -1.10325546  0.60439732
-0.10558415]
...
[ 0.3429808  0.00330087  0.14964075 ... -0.73518964 -0.68519336
-0.27575966]
[-0.84488505  0.1597866  -0.47073225 ... -0.24020459 -0.37110101
 1.17073215]
[-0.84488505 -0.8730192  0.04624525 ... -0.20212881 -0.47378505
-0.87137393]]
0      1
1      0
2      1
3      0
4      1
..
763    0
764    0
765    0
766    1
767    0
Name: Outcome, Length: 768, dtype: int64
```

### Train Test Split

```
X_train, X_test, Y_train, Y_test = train_test_split(X,Y, test_size = 0.2, stratify=Y, random_state=2)

print(X.shape, X_train.shape, X_test.shape)

(768, 8) (614, 8) (154, 8)
```

### Training the model

```
classifier = svm.SVC(kernel='linear')

# training the support vector machine classifier
classifier.fit(X_train, Y_train)
```

```
▼      SVC
SVC(kernel='linear')
```

### Model Evaluation

#### Accuracy score

```
# accuracy score on the training data
X_train_prediction = classifier.predict(X_train)
training_data_accuracy = accuracy_score(X_train_prediction, Y_train)

print('accuracy score of the training data:', training_data_accuracy)

accuracy score of the training data: 0.7866449511400652

# accuracy score on the test data
X_test_prediction = classifier.predict(X_test)
test_data_accuracy = accuracy_score(X_test_prediction, Y_test)

print('accuracy score of the test data:', test_data_accuracy)

accuracy score of the test data: 0.7727272727272727
```

### Making the predictive system

```
input_data = (4,110,92,0,0,37.6,0.191,30)

# changint the input_data to the numpy array
input_data_as_numpy_array = np.asarray(input_data)

# reshape the array as we are predicting for one instance
input_data_reshaped = input_data_as_numpy_array.reshape(1,-1)

# standarized the input_data
std_data = scaler.transform(input_data_reshaped)
print(std_data)
```

```
prediction = classifier.predict(std_data)
print(prediction)

if (prediction[0] == 0):
    print('the person is not diabetic')
else:
    print('the person is diabetic')

[[ 0.04601433 -0.34096773  1.18359575 -1.28821221 -0.69289057  0.71168975
 -0.84827977 -0.27575966]]
[0]
the person is not diabetic
/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but StandardScaler was fitted
  warnings.warn(
```